

Cortex Code (CoCo) — Developer Cheatsheet

[!IMPORTANT] Community cheatsheet — not official Snowflake documentation. For the authoritative reference, see [Cortex Code CLI docs](#).

Table of Contents

- [Quick Start](#)
- [Start Here By Role](#)
- [The 7 Trigger Characters](#)
- [Slash Commands You'll Actually Use](#)
- [Keyboard Shortcuts That Matter](#)
- [Working with Your Codebase](#)
- [Snowflake Data Shortcuts](#)
- [Agent Modes](#)
- [Skills](#)
- [Scheduling](#)
- [Memory Across Sessions](#)
- [Environment Variables](#)
- [Settings Worth Knowing](#)
- [CLI Subcommands](#)
- [MCP Integration](#)
- [Hooks](#)
- [Tips & Gotchas](#)
- [References](#)

Quick Start

Install on macOS/Linux:

```
# run in terminal
curl -sL https://sfc-repo.snowflakecomputing.com/cortex/install.sh
| bash
```

Enable shell completions (one-time setup):

```
# run in terminal
cortex completion bash > ~/.bash_completion.d/cortex # bash
cortex completion zsh > ~/.zsh/completions/_cortex # zsh
cortex completion fish > ~/.config/fish/completions/cortex.fish #
```

```
fish
exec $SHELL
```

Common launch options:

```
# run in terminal
cortex #
launch in current directory
cortex -c my_connection -w /path/to/project #
set connection and working dir
cortex --continue #
resume last session
cortex --resume <session_id> #
resume specific session by ID
cortex -p "list all Python files" --output-format stream-json #
one-off prompt (scripting)
cortex -f request.txt #
read prompt from file and exit
cortex -m claude-opus-4-6 -p "summarize this repo" #
specify model
```

Connections are defined in ~/.snowflake/connections.toml.
Recommended: OAUTH_AUTHORIZATION_CODE (browser-based OAuth — no stored password, tokens refreshed automatically):

```
[my_connection]
account = "<ACCOUNT_IDENTIFIER>"
user = "<USERNAME>"
authenticator = "OAUTH_AUTHORIZATION_CODE"
client_store_temporary_credential = true
role = "<ROLE>"
database = "<DATABASE>"
```

Start Here By Role

If you're a...	Focus on
DB / Data Developer	#TABLE, /sql, cortex search table-details, Semantic Views, SQL timeout
App Developer (Streamlit, Native Apps)	@{file}, /diff, /worktree, \$snowflake-apps, \$deploy-to-spcs
Data / ML Engineer	/lineage, \$machine-learning, \$cortex-agent, \$dynamic-tables, fdbt
Platform / Infra	/plan, /team, /bg, Hooks, MCP, Scheduling, Memory

The 7 Trigger Characters

Trigger	Name	What it does	Example
/	Slash command	Invoke a slash command	/sql SELECT current_user()
!	Bash shell	Run a terminal command inline	! git status
@	File reference	Reference a file by path (no inject)	@src/app.py review this
@{	File inject	Inject full file contents into prompt	@{schema.sql} write tests
#	Table reference	Autocomplete table name, load schema	#DB.PUBLIC.ORDERS summarize
\$	Skill invoke	Invoke a skill by name	\$semantic-view create a view for sales
%	Agent mention	Mention a Cortex Agent	%my_sales_agent what was Q1 revenue?

[!TIP] @{file} is the most useful trigger for developers: inject a config, schema, or source file into context without manually copying content.

Slash Commands You'll Actually Use

Execution & modes:

Command	What it does
/plan	Enter plan mode — CoCo presents a plan before making changes
/plan-off	Disable plan mode
/bypass	Auto-approve all tool calls
/bypass-off	Disable bypass mode
/team	Enable parallel teammates mode
/bg	Launch a background agent, keep chatting

Command	What it does
/agents	View and manage running subagents
/sandbox	Manage sandbox settings (on / off / status)

SQL & data:

Command	What it does
/sql <query>	Execute SQL inline (--limit N for row cap)
/sql-readonly	Toggle SQL write protection
/table	Open interactive table viewer
/connections	Manage Snowflake connections
/doctor	Diagnose connection issues

Code & git:

Command	What it does
/diff	Review git changes fullscreen (--staged for staged)
/worktree <create\ list\ switch\ delete>	Manage git worktrees
/dbt	dbt operations and lineage

Skills & config: /skill /rules /hooks /mcp /settings /permissions

Utilities: /copy /share /secrets /compact /rewind /new /resume /fork /index /tgrep /tasks /rename /update

Keyboard Shortcuts That Matter

Shortcut	Action
Ctrl+P	Toggle compact/expanded mode
Ctrl+O	Open full transcript viewer
Ctrl+T	Open table viewer
Ctrl+B	View background bash processes
Ctrl+D	Open todo/task viewer
Ctrl+C	Interrupt / cancel
Shift+Tab	Cycle mode (Confirm → Plan → Bypass)
Ctrl+R	History search (reverse)
Ctrl+J	Insert newline in input (for multiline prompts)
Esc Esc	Clear entire input

Shortcut	Action
?	Toggle help overlay
Ctrl+A	Move to start of input
Ctrl+E	Move to end of input

Working with Your Codebase

Trigger	Example	What happens
@file	@src/auth.py what does this do?	References the path (no content injection)
@{file}	@{schema.sql} write tests for this	Injects full file contents into context
@{f1} @{f2}	@{schema.sql} @{models.py} do these match?	Inject multiple files at once

For semantic code search: /tgrep status to check, /index --rebuild after major changes.

Ask naturally — CoCo uses tgrep internally: "find where authentication errors are handled"

Snowflake Data Shortcuts

Reference a table to load its schema instantly:

```
#SNOWFLAKE_SAMPLE_DATA.TPCH_SF1.ORDERS how many orders this year?
```

Run SQL (use Ctrl+J for newlines in multiline queries):

```
/sql SELECT COUNT(*), SUM(O_TOTALPRICE) FROM
SNOWFLAKE_SAMPLE_DATA.TPCH_SF1.ORDERS
    WHERE O_ORDERDATE >= '1995-01-01'
/sql SELECT O_ORDERSTATUS, COUNT(*) AS cnt
    FROM SNOWFLAKE_SAMPLE_DATA.TPCH_SF1.ORDERS GROUP BY 1
/sql SELECT * FROM SNOWFLAKE_SAMPLE_DATA.TPCH_SF1.CUSTOMER --limit
10
```

Discover objects and check docs before writing queries:

```
# run in terminal
cortex search object "TPCH orders"
cortex search table-details
"SNOWFLAKE_SAMPLE_DATA.TPCH_SF1.ORDERS"
cortex search docs "dynamic tables target lag"
```

```
cortex lineage SNOWFLAKE_SAMPLE_DATA.TPCH_SF1.ORDERS --direction
downstream --tree
```

Or from within CoCo: `!cortex search table-details`
`"SNOWFLAKE_SAMPLE_DATA.TPCH_SF1.ORDERS"`

SNOWFLAKE_SAMPLE_DATA is in every Snowflake account (no setup).
Other schemas: TPCDS_SF10TCL (retail), WEATHER.

Agent Modes

Mode	How to activate	When to use
Plan	<code>/plan</code>	Any multi-file or risky change — CoCo drafts a plan before touching anything
Bypass	<code>/bypass</code>	Auto-approve all tool calls (careful on prod)
Team	<code>/team</code>	Spawn parallel agents for large tasks
Background	<code>/bg</code>	Fire off a task, keep chatting; results arrive as a notification

Shift+Tab cycles through all modes: Confirm → Plan → Bypass.

Git worktrees for parallel safe work: `/worktree create <branch>`
gives each agent its own branch.

Skills

Invoke with `$` prefix in the CoCo prompt:
`$cortex-agent build a Q&A agent over my warehouse`

Manage:

```
/skill          # browse installed skills
/skill list     # list all
```

```
# run in terminal
cortex skill list
cortex skill find "<query>"
cortex skill add <name>
```

Custom skills: place a .md file in .cortex/skills/ (project) or ~/.snowflake/cortex/skills/ (global).

Bundled skills:

Skill	What it does
<code>\$cortex-agent</code>	Build, debug, deploy Cortex Agents
<code>\$semantic-view</code>	Create/manage semantic views for Cortex Analyst
<code>\$cortex-ai-functions</code>	AI_CLASSIFY, AI_SENTIMENT, embeddings, OCR at scale
<code>\$machine-learning</code>	Train models, use Model Registry, run ML jobs
<code>\$dynamic-tables</code>	Build incremental pipelines with Dynamic Tables
<code>\$snowpark-python</code>	Snowpark Python UDFs, UDAFs, stored procedures
<code>\$iceberg</code>	Iceberg tables: catalog integrations, external volumes
<code>\$snowflake-apps</code>	Scaffold and deploy Snowflake App Runtime apps
<code>\$developing-with-streamlit</code>	Streamlit in Snowflake: theming, custom components
<code>\$snowflake-notebooks</code>	Workspace notebooks: Snowpark Python + SQL cells
<code>\$deploy-to-spcs</code>	Deploy containers to Snowpark Container Services
<code>\$workload-performance-analysis</code>	SQL query performance: spilling, partition pruning
<code>\$lineage</code>	Trace upstream/downstream column-level lineage
<code>\$warehouse</code>	Analyze and right-size warehouse costs

Full list (30+): [Bundled Skills](#)

Scheduling

<code>/loop</code>	<code># interactive scheduler</code>
<code>/loop "every day at 9am"</code>	<code># natural language</code>
<code>/loop "0 9 * * 1-5"</code>	<code># standard cron (Mon–Fri 9am)</code>
<code>/automation</code>	<code># schedule a recurring agent task</code>

[!NOTE] Scheduled tasks are session-scoped and expire after 3 days. Max 50 per session.

Memory Across Sessions

Enable once, then memories persist across every CoCo session.

```
# run in terminal – enable memory (pick one)
export CORTEX_ENABLE_MEMORY=1          # environment variable
# or set in ~/.snowflake/cortex/settings.json: "enableMemory":
true
```

remember — save something CoCo should know in every future session:

```
cortex memory remember "use conventional commits: feat:, fix:,
docs:, chore:, refactor:, test:"
cortex memory remember "production db is read-only, use role
PROD_READ for queries"
cortex memory remember "default warehouse is COMPUTE_WH_M, only
use XL for full-table scans"
cortex memory remember "prefer async/await over .then() in this
codebase"
```

recall — find a specific memory by topic:

```
cortex memory recall "commit format"
cortex memory recall "production access"
cortex memory recall "warehouse sizing"
```

list / drop — manage stored memories:

```
cortex memory list          # show all memories with IDs
cortex memory drop <id>    # remove a specific memory by ID
```

[!TIP] Natural language works too:

- Say “remember that we deploy to us-east-1” and CoCo stores it.
- Say “what’s our commit convention?” and CoCo recalls it from memory without a command.

Environment Variables

Variable	Effect
CORTEX_ENABLE_MEMORY	Enable cross-session memory

Variable	Effect
	(alternative to settings)
COCO_DISABLE_CRON	Disable /loop and all scheduling tools
COCO_DISABLE_ROUTINES	Disable routines
CORTEX_CODE_STREAMING	Enable code streaming
CORTEX_AGENT_USE_LOCAL_ORCHESTRATOR	Use local orchestrator (developer/debug mode)
CTX_STEP_ENFORCEMENT	Enforce ctx step tracking discipline
CORTEX_BROWSER_HEADLESS	Run browser automation without a visible window
CORTEX_CODE_ENABLE_SNOWFLAKE_MANAGED_MCP_SERVERS	Enable Snowflake-managed MCP servers
CORTEX_DISABLE_TODO_TOOL	Disable the todo/task tool

Set inline for a single session:

```
# run in terminal
CORTEX_ENABLE_MEMORY=1 CORTEX_BROWSER_HEADLESS=1 cortex
```

Settings Worth Knowing

Configure via /settings or edit ~/.snowflake/cortex/settings.json.

Key	Default	What it does
agentMode	standard	standard = balanced; code = optimized for code-heavy tasks (higher tool call budget)

Key	Default	What it does
autoAcceptPlans	false	Skip the “approve plan?” prompt
enableMemory	false	Enable cross-session memory
tgrepEnabled	true	Enable semantic code search
sqlDefaultTimeoutSeconds	180	Max wait for SQL queries
diffDisplayMode	unified	unified (git-style) or side_by_side
bashDefaultTimeoutMs	180000	Timeout for shell commands
mcpWait	false	Wait for all MCP servers before starting (useful in CI)
disableCron	false	Disable /loop scheduling
autoUpdate	true	false to disable automatic version updates

Minimal ~/.snowflake/cortex/settings.json to start from:

```
{
  "enableMemory": true,
  "autoAcceptPlans": false,
  "autoUpdate": false,
  "sqlDefaultTimeoutSeconds": 300
}
```

CLI Subcommands

Run outside the interactive session — useful in scripts, CI, and automation.

Command	What it does
cortex search object "<query>"	Find tables, views, schemas by name or description
cortex search table-details "DB.SCHEMA.TABLE"	Full column metadata for a table
cortex search docs "<query>"	Search Snowflake product docs Search Marketplace listings

Command	What it does
cortex search marketplace "<query>"	
cortex lineage DB.SCHEMA.TABLE --direction both --tree	Trace upstream and downstream lineage
cortex connections list	List available connections
cortex connections set <name>	Switch active connection
cortex semantic-views list	List semantic views in the account
cortex worktree create <branch>	Create a git worktree (also available via /worktree)
cortex skill list	List installed skills
cortex memory remember "<text>"	Store a memory (requires enableMemory: true)
cortex update	Update to the latest version

For the full command list: [CLI Reference](#)

MCP Integration

Config file: ~/.snowflake/cortex/mcp.json | Manage: cortex mcp add | list | remove | In-session: /mcp

Tool naming inside CoCo: mcp__<server>__<tool>
(e.g. mcp__github__create_pull_request)

Supported transports: http, sse, stdio

Add a server (example: GitHub):

```
# run in terminal
cortex mcp add
# ↳ transport=stdio, command=npx, args=-y @modelcontextprotocol/
server-github
# ↳ env: GITHUB_PERSONAL_ACCESS_TOKEN=<your-token>
```

Other popular servers (all via npx -y @modelcontextprotocol/server-<name>): filesystem, postgres, brave-search, slack

Set mcpWait: true in settings if MCP tools must be ready before the first prompt (CI/automation).

```
# run in terminal
cortex mcp list          # show configured servers
cortex mcp remove <server_name> # remove a server
```

For setup details: [MCP docs](#)

Hooks

Hooks fire shell scripts in response to lifecycle events. Configure in `~/.snowflake/cortex/hooks/`.

Event	When it fires
UserPromptSubmit	User submits a prompt
PreToolUse	Before any tool call
PostToolUse	After any tool call
PermissionRequest	When a permission prompt appears
Stop	Agent stops
SessionStart	Session begins
SessionEnd	Session ends
PreCompact	Before conversation compaction

View and test configured hooks:

`/hooks`

Tips & Gotchas

- **/plan before anything risky.** CoCo reads the codebase, drafts a plan, and waits for your approval. For multi-file changes this is the difference between a clean refactor and a mess.
- **@file does NOT inject content — @{file} does.** `@src/auth.py` review this just references the path; CoCo may or may not read it. `@{src/auth.py}` review this injects the full contents every time. If CoCo seems to “not know” a file, switch to `@{`.
- **#TABLE loads schema automatically.** Type `#` then the table name — CoCo fetches column definitions before writing any SQL. No need to describe the table.
- **/bg for long jobs.** Fire off a background agent to run a big refactor or analysis while you keep chatting. Results arrive as a notification.
- **cortex search docs is semantic.** “how do I set target lag for dynamic tables” finds the right page faster than a Google search.
- **Secrets stay out of chat.** Use `/secrets` to store API keys and tokens. CoCo injects them at runtime without exposing them in the session transcript.

References

- [Cortex Code CLI](#)
- [CLI Reference \(slash commands, batch mode\)](#)
- [Keyboard Shortcuts](#)

- [Bundled Skills](#)
- [Settings Reference](#)
- [Extensibility \(hooks, plugins, MCP\)](#)
- [Changelog](#)